

Package ‘nnet’

August 30, 2015

Priority recommended

Version 7.3-11

Date 2015-08-29

Depends R (>= 2.14.0), stats, utils

Suggests MASS

Description Software for feed-forward neural networks with a single hidden layer, and for multinomial log-linear models.

Title Feed-Forward Neural Networks and Multinomial Log-Linear Models

ByteCompile yes

License GPL-2 | GPL-3

URL <http://www.stats.ox.ac.uk/pub/MASS4/>

NeedsCompilation yes

Author Brian Ripley [aut, cre, cph],
William Venables [cph]

Maintainer Brian Ripley <ripley@stats.ox.ac.uk>

Repository CRAN

Date/Publication 2015-08-30 08:59:40

R topics documented:

| | |
|------------------------|---|
| class.ind | 2 |
| multinom | 2 |
| nnet | 4 |
| nnetHess | 7 |
| predict.nnet | 8 |
| which.is.max | 9 |

| | |
|--------------|-----------|
| Index | 11 |
|--------------|-----------|

| | |
|-----------|---|
| class.ind | <i>Generates Class Indicator Matrix from a Factor</i> |
|-----------|---|

Description

Generates a class indicator function from a given factor.

Usage

```
class.ind(cl)
```

Arguments

cl factor or vector of classes for cases.

Value

a matrix which is zero except for the column corresponding to the class.

References

Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. Fourth edition. Springer.

Examples

```
# The function is currently defined as
class.ind <- function(cl)
{
  n <- length(cl)
  cl <- as.factor(cl)
  x <- matrix(0, n, length(levels(cl)) )
  x[(1:n) + n*(unclass(cl)-1)] <- 1
  dimnames(x) <- list(names(cl), levels(cl))
  x
}
```

| | |
|----------|--|
| multinom | <i>Fit Multinomial Log-linear Models</i> |
|----------|--|

Description

Fits multinomial log-linear models via neural networks.

Usage

```
multinom(formula, data, weights, subset, na.action,
          contrasts = NULL, Hess = FALSE, summ = 0, censored = FALSE,
          model = FALSE, ...)
```

Arguments

| | |
|-----------|---|
| formula | a formula expression as for regression models, of the form <code>response ~ predictors</code> . The response should be a factor or a matrix with K columns, which will be interpreted as counts for each of K classes. A log-linear model is fitted, with coefficients zero for the first class. An offset can be included: it should be a numeric matrix with K columns if the response is either a matrix with K columns or a factor with $K \geq 2$ classes, or a numeric vector for a response factor with 2 levels. See the documentation of <code>formula()</code> for other details. |
| data | an optional data frame in which to interpret the variables occurring in formula. |
| weights | optional case weights in fitting. |
| subset | expression saying which subset of the rows of the data should be used in the fit. All observations are included by default. |
| na.action | a function to filter missing data. |
| contrasts | a list of contrasts to be used for some or all of the factors appearing as variables in the model formula. |
| Hess | logical for whether the Hessian (the observed/expected information matrix) should be returned. |
| summ | integer; if non-zero summarize by deleting duplicate rows and adjust weights. Methods 1 and 2 differ in speed (2 uses C); method 3 also combines rows with the same X and different Y, which changes the baseline for the deviance. |
| censored | If Y is a matrix with K columns, interpret the entries as one for possible classes, zero for impossible classes, rather than as counts. |
| model | logical. If true, the model frame is saved as component <code>model</code> of the returned object. |
| ... | additional arguments for <code>nnet</code> |

Details

`multinom` calls `nnet`. The variables on the rhs of the formula should be roughly scaled to [0,1] or the fit will be slow or may not converge at all.

Value

A `nnet` object with additional components:

| | |
|----------|--|
| deviance | the residual deviance, compared to the full saturated model (that explains individual observations exactly). Also, minus twice log-likelihood. |
| edf | the (effective) number of degrees of freedom used by the model |
| AIC | the AIC for this fit. |
| Hessian | (if <code>Hess</code> is true). |
| model | (if <code>model</code> is true). |

References

Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. Fourth edition. Springer.

See Also[nnet](#)**Examples**

```
options(contrasts = c("contr.treatment", "contr.poly"))
library(MASS)
example(birthwt)
(bwt.mu <- multinom(low ~ ., bwt))
```

nnet*Fit Neural Networks*

Description

Fit single-hidden-layer neural network, possibly with skip-layer connections.

Usage

```
nnet(x, ...)

## S3 method for class 'formula'
nnet(formula, data, weights, ...,
      subset, na.action, contrasts = NULL)

## Default S3 method:
nnet(x, y, weights, size, Wts, mask,
      linout = FALSE, entropy = FALSE, softmax = FALSE,
      censored = FALSE, skip = FALSE, rang = 0.7, decay = 0,
      maxit = 100, Hess = FALSE, trace = TRUE, MaxNWts = 1000,
      abstol = 1.0e-4, reltol = 1.0e-8, ...)
```

Arguments

| | |
|---------|--|
| formula | A formula of the form <code>class ~ x1 + x2 + ...</code> |
| x | matrix or data frame of x values for examples. |
| y | matrix or data frame of target values for examples. |
| weights | (case) weights for each example – if missing defaults to 1. |
| size | number of units in the hidden layer. Can be zero if there are skip-layer units. |
| data | Data frame from which variables specified in formula are preferentially to be taken. |
| subset | An index vector specifying the cases to be used in the training sample. (NOTE: If given, this argument must be named.) |

| | |
|------------------------|---|
| <code>na.action</code> | A function to specify the action to be taken if NAs are found. The default action is for the procedure to fail. An alternative is <code>na.omit</code> , which leads to rejection of cases with missing values on any required variable. (NOTE: If given, this argument must be named.) |
| <code>contrasts</code> | a list of contrasts to be used for some or all of the factors appearing as variables in the model formula. |
| <code>Wts</code> | initial parameter vector. If missing chosen at random. |
| <code>mask</code> | logical vector indicating which parameters should be optimized (default all). |
| <code>linout</code> | switch for linear output units. Default logistic output units. |
| <code>entropy</code> | switch for entropy (= maximum conditional likelihood) fitting. Default by least-squares. |
| <code>softmax</code> | switch for softmax (log-linear model) and maximum conditional likelihood fitting. <code>linout</code> , <code>entropy</code> , <code>softmax</code> and <code>censored</code> are mutually exclusive. |
| <code>censored</code> | A variant on <code>softmax</code> , in which non-zero targets mean possible classes. Thus for <code>softmax</code> a row of (0, 1, 1) means one example each of classes 2 and 3, but for <code>censored</code> it means one example whose class is only known to be 2 or 3. |
| <code>skip</code> | switch to add skip-layer connections from input to output. |
| <code>rang</code> | Initial random weights on [-rang, rang]. Value about 0.5 unless the inputs are large, in which case it should be chosen so that <code>rang * max(x)</code> is about 1. |
| <code>decay</code> | parameter for weight decay. Default 0. |
| <code>maxit</code> | maximum number of iterations. Default 100. |
| <code>Hess</code> | If true, the Hessian of the measure of fit at the best set of weights found is returned as component Hessian. |
| <code>trace</code> | switch for tracing optimization. Default TRUE. |
| <code>MaxNWts</code> | The maximum allowable number of weights. There is no intrinsic limit in the code, but increasing <code>MaxNWts</code> will probably allow fits that are very slow and time-consuming. |
| <code>abstol</code> | Stop if the fit criterion falls below <code>abstol</code> , indicating an essentially perfect fit. |
| <code>reltol</code> | Stop if the optimizer is unable to reduce the fit criterion by a factor of at least <code>1 - reltol</code> . |
| <code>...</code> | arguments passed to or from other methods. |

Details

If the response in `formula` is a factor, an appropriate classification network is constructed; this has one output and entropy fit if the number of levels is two, and a number of outputs equal to the number of classes and a softmax output stage for more levels. If the response is not a factor, it is passed on unchanged to `nnet.default`.

Optimization is done via the BFGS method of [optim](#).

Value

object of class "nnet" or "nnet.formula". Mostly internal structure, but has components

| | |
|----------------------------|---|
| <code>wt</code> | the best set of weights found |
| <code>value</code> | value of fitting criterion plus weight decay term. |
| <code>fitted.values</code> | the fitted values for the training data. |
| <code>residuals</code> | the residuals for the training data. |
| <code>convergence</code> | 1 if the maximum number of iterations was reached, otherwise 0. |

References

Ripley, B. D. (1996) *Pattern Recognition and Neural Networks*. Cambridge.

Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. Fourth edition. Springer.

See Also

[predict.nnet](#), [nnetHess](#)

Examples

```
# use half the iris data
ir <- rbind(iris3[,1],iris3[,2],iris3[,3])
targets <- class.ind( c(rep("s", 50), rep("c", 50), rep("v", 50)) )
samp <- c(sample(1:50,25), sample(51:100,25), sample(101:150,25))
ir1 <- nnet(ir[samp,], targets[samp,], size = 2, rang = 0.1,
            decay = 5e-4, maxit = 200)
test.cl <- function(true, pred) {
  true <- max.col(true)
  cres <- max.col(pred)
  table(true, cres)
}
test.cl(targets[-samp,], predict(ir1, ir[-samp,]))

# or
ird <- data.frame(rbind(iris3[,1], iris3[,2], iris3[,3]),
                 species = factor(c(rep("s",50), rep("c", 50), rep("v", 50))))
ir.nn2 <- nnet(species ~ ., data = ird, subset = samp, size = 2, rang = 0.1,
              decay = 5e-4, maxit = 200)
table(ird$species[-samp], predict(ir.nn2, ird[-samp,], type = "class"))
```

nnetHess*Evaluates Hessian for a Neural Network*

Description

Evaluates the Hessian (matrix of second derivatives) of the specified neural network. Normally called via argument Hess=TRUE to `nnet` or via `vcov.multinom`.

Usage

```
nnetHess(net, x, y, weights)
```

Arguments

| | |
|----------------------|--|
| <code>net</code> | object of class <code>nnet</code> as returned by <code>nnet</code> . |
| <code>x</code> | training data. |
| <code>y</code> | classes for training data. |
| <code>weights</code> | the (case) weights used in the <code>nnet</code> fit. |

Value

square symmetric matrix of the Hessian evaluated at the weights stored in the `net`.

References

Ripley, B. D. (1996) *Pattern Recognition and Neural Networks*. Cambridge.
Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. Fourth edition. Springer.

See Also

[nnet](#), [predict.nnet](#)

Examples

```
# use half the iris data
ir <- rbind(iris3[,1], iris3[,2], iris3[,3])
targets <- matrix(c(rep(c(1,0,0),50), rep(c(0,1,0),50), rep(c(0,0,1),50)),
  150, 3, byrow=TRUE)
samp <- c(sample(1:50,25), sample(51:100,25), sample(101:150,25))
ir1 <- nnet(ir[samp,], targets[samp,], size=2, rang=0.1, decay=5e-4, maxit=200)
eigen(nnetHess(ir1, ir[samp,], targets[samp,]), TRUE)$values
```

predict.nnet

Predict New Examples by a Trained Neural Net

Description

Predict new examples by a trained neural net.

Usage

```
## S3 method for class 'nnet'
predict(object, newdata, type = c("raw", "class"), ...)
```

Arguments

| | |
|---------|--|
| object | an object of class nnet as returned by nnet. |
| newdata | matrix or data frame of test examples. A vector is considered to be a row vector comprising a single case. |
| type | Type of output |
| ... | arguments passed to or from other methods. |

Details

This function is a method for the generic function `predict()` for class "nnet". It can be invoked by calling `predict(x)` for an object `x` of the appropriate class, or directly by calling `predict.nnet(x)` regardless of the class of the object.

Value

If `type = "raw"`, the matrix of values returned by the trained network; if `type = "class"`, the corresponding class (which is probably only useful if the net was generated by `nnet.formula`).

References

Ripley, B. D. (1996) *Pattern Recognition and Neural Networks*. Cambridge.

Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. Fourth edition. Springer.

See Also

[nnet](#), [which.is.max](#)

Examples

```
# use half the iris data
ir <- rbind(iris3[,1], iris3[,2], iris3[,3])
targets <- class.ind( c(rep("s", 50), rep("c", 50), rep("v", 50)) )
samp <- c(sample(1:50,25), sample(51:100,25), sample(101:150,25))
ir1 <- nnet(ir[samp,], targets[samp,],size = 2, rang = 0.1,
            decay = 5e-4, maxit = 200)
test.cl <- function(true, pred){
  true <- max.col(true)
  cres <- max.col(pred)
  table(true, cres)
}
test.cl(targets[-samp,], predict(ir1, ir[-samp,]))

# or
ird <- data.frame(rbind(iris3[,1], iris3[,2], iris3[,3]),
                  species = factor(c(rep("s",50), rep("c", 50), rep("v", 50))))
ir.nn2 <- nnet(species ~ ., data = ird, subset = samp, size = 2, rang = 0.1,
              decay = 5e-4, maxit = 200)
table(ird$species[-samp], predict(ir.nn2, ird[-samp,], type = "class"))
```

which.is.max

Find Maximum Position in Vector

Description

Find the maximum position in a vector, breaking ties at random.

Usage

```
which.is.max(x)
```

Arguments

x a vector

Details

Ties are broken at random.

Value

index of a maximal value.

References

Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. Fourth edition. Springer.

See Also

[max.col](#), [which.max](#) which takes the first of ties.

Examples

```
## Not run: ## this is incomplete
pred <- predict(nnet, test)
table(true, apply(pred, 1, which.is.max))

## End(Not run)
```

Index

*Topic **models**

multinom, [2](#)

*Topic **neural**

class.ind, [2](#)

multinom, [2](#)

nnet, [4](#)

nnetHess, [7](#)

predict.nnet, [8](#)

*Topic **utilities**

class.ind, [2](#)

which.is.max, [9](#)

add.net (nnet), [4](#)

add1.multinom (multinom), [2](#)

anova.multinom (multinom), [2](#)

class.ind, [2](#)

coef.multinom (multinom), [2](#)

coef.nnet (nnet), [4](#)

drop1.multinom (multinom), [2](#)

eval.nn (nnet), [4](#)

extractAIC.multinom (multinom), [2](#)

formula, [3](#)

logLik.multinom (multinom), [2](#)

max.col, [10](#)

model.frame.multinom (multinom), [2](#)

multinom, [2](#)

nnet, [3](#), [4](#), [4](#), [7](#), [8](#)

nnetHess, [6](#), [7](#)

norm.net (nnet), [4](#)

optim, [5](#)

predict.multinom (multinom), [2](#)

predict.nnet, [6](#), [7](#), [8](#)

print.multinom (multinom), [2](#)

print.nnet (nnet), [4](#)

print.summary.multinom (multinom), [2](#)

print.summary.nnet (nnet), [4](#)

summary.multinom (multinom), [2](#)

summary.nnet (nnet), [4](#)

vcov.multinom (multinom), [2](#)

which.is.max, [8](#), [9](#)

which.max, [10](#)